

Bibliography for seminar

This file only contains paper references. References for image are not included.

References

- [18] “Safety and Trustworthiness of Deep Neural Networks: A Survey”. In: *CoRR* abs/1812.08342 (2018). Withdrawn. arXiv: 1812.08342. URL: <http://arxiv.org/abs/1812.08342>.
- [Ans+19] Alessio Ansuini et al. “Intrinsic dimension of data representations in deep neural networks”. In: *CoRR* abs/1905.12784 (2019). arXiv: 1905.12784. URL: <http://arxiv.org/abs/1905.12784>.
- [BCB16] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. *Neural Machine Translation by Jointly Learning to Align and Translate*. 2016. arXiv: 1409.0473 [cs.CL].
- [BD07] Alan J. Bray and David S. Dean. “Statistics of Critical Points of Gaussian Fields on Large-Dimensional Spaces”. In: *Physical Review Letters* 98.15 (Apr. 2007). ISSN: 1079-7114. DOI: 10.1103/physrevlett.98.150201. URL: <http://dx.doi.org/10.1103/PhysRevLett.98.150201>.
- [Bel+19] Mikhail Belkin et al. *Reconciling modern machine learning practice and the bias-variance trade-off*. 2019. arXiv: 1812.11118 [stat.ML].
- [Cyb89] G. Cybenko. “Approximation by superpositions of a sigmoidal function”. In: *Mathematics of Control, Signals and Systems* 2.4 (Dec. 1989), pp. 303–314. ISSN: 1435-568X. DOI: 10.1007/BF02551274. URL: <https://doi.org/10.1007/BF02551274>.
- [Dau+14] Yann N Dauphin et al. “Identifying and attacking the saddle point problem in high-dimensional non-convex optimization”. In: *Advances in Neural Information Processing Systems*. Ed. by Z. Ghahramani et al. Vol. 27. Curran Associates, Inc., 2014. URL: <https://proceedings.neurips.cc/paper/2014/file/17e23e50bedc63b4095e3d8204ce063b-Paper.pdf>.
- [Du+18] Simon S. Du et al. “Gradient Descent Finds Global Minima of Deep Neural Networks”. In: *CoRR* abs/1811.03804 (2018). arXiv: 1811.03804. URL: <http://arxiv.org/abs/1811.03804>.
- [FMN13] Charles Fefferman, Sanjoy Mitter, and Hariharan Narayanan. *Testing the Manifold Hypothesis*. 2013. arXiv: 1310.0425 [math.ST].
- [Geh+18] Timon Gehr et al. “AI2: Safety and Robustness Certification of Neural Networks with Abstract Interpretation”. In: *2018 IEEE Symposium on Security and Privacy (SP)*. 2018, pp. 3–18. DOI: 10.1109/SP.2018.00058.

- [GL20] Shuyue Guan and Murray Loew. “Analysis of Generalizability of Deep Neural Networks Based on the Complexity of Decision Boundary”. In: *2020 19th IEEE International Conference on Machine Learning and Applications (ICMLA)* (Dec. 2020). DOI: 10.1109/icmla51294.2020.00025. URL: <http://dx.doi.org/10.1109/ICMLA51294.2020.00025>.
- [Gui+18] Riccardo Guidotti et al. “A Survey of Methods for Explaining Black Box Models”. In: *ACM Comput. Surv.* 51.5 (Aug. 2018). ISSN: 0360-0300. DOI: 10.1145/3236009. URL: <https://doi.org/10.1145/3236009>.
- [HJ15] Tamir Hazan and Tommi S. Jaakkola. “Steps Toward Deep Kernel Methods from Infinite Neural Networks”. In: *CoRR* abs/1508.05133 (2015). arXiv: 1508.05133. URL: <http://arxiv.org/abs/1508.05133>.
- [HR17] Michael Hauser and Asok Ray. “Principles of Riemannian Geometry in Neural Networks”. In: *Advances in Neural Information Processing Systems*. Ed. by I. Guyon et al. Vol. 30. Curran Associates, Inc., 2017. URL: <https://proceedings.neurips.cc/paper/2017/file/0ebcc77dc72360d0eb8e9504c78d38bd-Paper.pdf>.
- [HT21] Fengxiang He and Dacheng Tao. *Recent advances in deep learning theory*. 2021. arXiv: 2012.10931 [cs.LG].
- [HWT20] Fengxiang He, Bohan Wang, and Dacheng Tao. *Piecewise linear activations substantially shape the loss surfaces of neural networks*. 2020. arXiv: 2003.12236 [cs.LG].
- [Kat+17] Guy Katz et al. “Reluplex: An Efficient SMT Solver for Verifying Deep Neural Networks”. In: *CoRR* abs/1702.01135 (2017). arXiv: 1702.01135. URL: <http://arxiv.org/abs/1702.01135>.
- [KFY18] Jinhan Kim, Robert Feldt, and Shin Yoo. “Guiding Deep Learning System Testing using Surprise Adequacy”. In: *CoRR* abs/1808.08444 (2018). arXiv: 1808.08444. URL: <http://arxiv.org/abs/1808.08444>.
- [KKB20] Kenji Kawaguchi, Leslie Pack Kaelbling, and Yoshua Bengio. *Generalization in Deep Learning*. 2020. arXiv: 1710.05468 [stat.ML].
- [LBH15] Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. “Deep learning”. In: *Nature* 521.7553 (May 2015), pp. 436–444. ISSN: 1476-4687. DOI: 10.1038/nature14539. URL: <https://doi.org/10.1038/nature14539>.
- [LL20] Yulong Lu and Jianfeng Lu. *A Universal Approximation Theorem of Deep Neural Networks for Expressing Probability Distributions*. 2020. arXiv: 2004.08867 [cs.LG].

- [Ma+18] Lei Ma et al. “DeepGauge: Comprehensive and Multi-Granularity Testing Criteria for Gauging the Robustness of Deep Learning Systems”. In: *CoRR* abs/1803.07519 (2018). arXiv: 1803.07519. URL: <http://arxiv.org/abs/1803.07519>.
- [Ola14] Christopher Olah. *Neural Networks, Manifolds, and Topology*. <https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/>. Accessed: 2021-04-29. Apr. 2014.
- [Par+21] Sejun Park et al. “Minimum Width for Universal Approximation”. In: *International Conference on Learning Representations*. 2021. URL: <https://openreview.net/forum?id=0-XJwyoIF-k>.
- [Pei+17] Kexin Pei et al. “DeepXplore: Automated Whitebox Testing of Deep Learning Systems”. In: *CoRR* abs/1705.06640 (2017). arXiv: 1705.06640. URL: <http://arxiv.org/abs/1705.06640>.
- [Qin+20] Chongli Qin et al. *Training Generative Adversarial Networks by Solving Ordinary Differential Equations*. 2020. arXiv: 2010.15040 [stat.ML].
- [RHK18] Wenjie Ruan, Xiaowei Huang, and Marta Kwiatkowska. *Reachability Analysis of Deep Neural Networks with Provable Guarantees*. 2018. arXiv: 1805.02242 [cs.LG].
- [SHK18] Youcheng Sun, Xiaowei Huang, and Daniel Kroening. “Testing Deep Neural Networks”. In: *CoRR* abs/1803.04792 (2018). arXiv: 1803.04792. URL: <http://arxiv.org/abs/1803.04792>.
- [SZ06] Anton Maximilian Schäfer and Hans Georg Zimmermann. “Recurrent Neural Networks Are Universal Approximators”. In: *Proceedings of the 16th International Conference on Artificial Neural Networks - Volume Part I*. ICANN’06. Athens, Greece: Springer-Verlag, 2006, pp. 632–640. ISBN: 3540386254. DOI: 10.1007/11840817_66. URL: https://doi.org/10.1007/11840817_66.
- [UM21] Caterina Urban and Antoine Miné. *A Review of Formal Methods applied to Machine Learning*. 2021. arXiv: 2104.02466 [cs.PL].
- [Yos+15] Jason Yosinski et al. *Understanding Neural Networks Through Deep Visualization*. 2015. arXiv: 1506.06579 [cs.CV].
- [Zha+16] Chiyuan Zhang et al. “Understanding deep learning requires rethinking generalization”. In: *CoRR* abs/1611.03530 (2016). arXiv: 1611.03530. URL: <http://arxiv.org/abs/1611.03530>.
- [Zho20] Ding-Xuan Zhou. “Universality of deep convolutional neural networks”. In: *Applied and Computational Harmonic Analysis* 48.2 (2020), pp. 787–794. ISSN: 1063-5203. DOI: <https://doi.org/10.1016/j.acha.2019.06.004>. URL: <https://www.sciencedirect.com/science/article/pii/S1063520318302045>.
- [Zhu+19] He Zhu et al. “An Inductive Synthesis Framework for Verifiable Reinforcement Learning”. In: *CoRR* abs/1907.07273 (2019). arXiv: 1907.07273. URL: <http://arxiv.org/abs/1907.07273>.