Bibliography for seminar

References

[18]   "Safety and Trustworthiness of Deep Neural Networks: A Survey".
          In: *CoRR* abs/1812.08342 (2018). Withdrawn. arXiv: 1812.08342.
          URL: http://arxiv.org/abs/1812.08342.

[Ans+19] Alessio Ansuini et al. "Intrinsic dimension of data representations
          in deep neural networks". In: *CoRR* abs/1905.12784 (2019). arXiv:
          1905.12784. URL: http://arxiv.org/abs/1905.12784.

[BCB16]   Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. *Neural
          Machine Translation by Jointly Learning to Align and Translate.*
          2016. arXiv: 1409.0473 [cs.CL].

[BD07]   Alan J. Bray and David S. Dean. "Statistics of Critical Points of
          Gaussian Fields on Large-Dimensional Spaces". In: *Physical Re-
          view Letters* 98.15 (Apr. 2007). ISSN: 1079-7114. DOI: 10.1103/
          physrevlett.98.150201. URL: http://dx.doi.org/10.1103/
          PhysRevLett.98.150201.

[Bel+19] Mikhail Belkin et al. *Reconciling modern machine learning practice
          and the bias-variance trade-off.* 2019. arXiv: 1812.11118 [stat.ML].

[Cyb89]   G. Cybenko. "Approximation by superpositions of a sigmoidal func-
          tion". In: *Mathematics of Control, Signals and Systems* 2.4 (Dec.
          1989), pp. 303–314. ISSN: 1435-568X. DOI: 10.1007/BF02551274.
          URL: https://doi.org/10.1007/BF02551274.

[Dau+14] Yann N Dauphin et al. "Identifying and attacking the saddle point
          problem in high-dimensional non-convex optimization". In: *Advances
          in Neural Information Processing Systems.* Ed. by Z. Ghahramani et
          al. Vol. 27. Curran Associates, Inc., 2014. URL: https://proceedings.
          neurips.cc/paper/2014/file/17e23e50bedc63b4095e3d8204ce063b-
          Paper.pdf.

[Du+18]   Simon S. Du et al. "Gradient Descent Finds Global Minima of Deep
          Neural Networks". In: *CoRR* abs/1811.03804 (2018). arXiv: 1811.
          03804. URL: http://arxiv.org/abs/1811.03804.

[FMN13]   Charles Fefferman, Sanjoy Mitter, and Hariharan Narayanan. *Test-
          ing the Manifold Hypothesis.* 2013. arXiv: 1310.0425 [math.ST].

[Geh+18] Timon Gehr et al. "AI2: Safety and Robustness Certification of Neu-
          ral Networks with Abstract Interpretation". In: *2018 IEEE Sympo-
          sium on Security and Privacy (SP).* 2018, pp. 3–18. DOI: 10.1109/
          SP.2018.00058.

[GL20]   Shuyue Guan and Murray Loew. "Analysis of Generalizability of Deep Neural Networks Based on the Complexity of Decision Boundary". In: *2020 19th IEEE International Conference on Machine Learning and Applications (ICMLA)* (Dec. 2020). DOI: 10.1109/icmla51294.2020.00025. URL: http://dx.doi.org/10.1109/ICMLA51294.2020.00025.

[Gui+18]   Riccardo Guidotti et al. "A Survey of Methods for Explaining Black Box Models". In: *ACM Comput. Surv.* 51.5 (Aug. 2018). ISSN: 0360-0300. DOI: 10.1145/3236009. URL: https://doi.org/10.1145/3236009.

[HJ15]   Tamir Hazan and Tommi S. Jaakkola. "Steps Toward Deep Kernel Methods from Infinite Neural Networks". In: *CoRR* abs/1508.05133 (2015). arXiv: 1508.05133. URL: http://arxiv.org/abs/1508.05133.

[HR17]   Michael Hauser and Asok Ray. "Principles of Riemannian Geometry in Neural Networks". In: *Advances in Neural Information Processing Systems*. Ed. by I. Guyon et al. Vol. 30. Curran Associates, Inc., 2017. URL: https://proceedings.neurips.cc/paper/2017/file/0ebcc77dc72360d0eb8e9504c78d38bd-Paper.pdf.

[HT21]   Fengxiang He and Dacheng Tao. *Recent advances in deep learning theory*. 2021. arXiv: 2012.10931 [cs.LG].

[HWT20]   Fengxiang He, Bohan Wang, and Dacheng Tao. *Piecewise linear activations substantially shape the loss surfaces of neural networks*. 2020. arXiv: 2003.12236 [cs.LG].

[Kat+17]   Guy Katz et al. "Reluplex: An Efficient SMT Solver for Verifying Deep Neural Networks". In: *CoRR* abs/1702.01135 (2017). arXiv: 1702.01135. URL: http://arxiv.org/abs/1702.01135.

[KFY18]   Jinhan Kim, Robert Feldt, and Shin Yoo. "Guiding Deep Learning System Testing using Surprise Adequacy". In: *CoRR* abs/1808.08444 (2018). arXiv: 1808.08444. URL: http://arxiv.org/abs/1808.08444.

[KKB20]   Kenji Kawaguchi, Leslie Pack Kaelbling, and Yoshua Bengio. *Generalization in Deep Learning*. 2020. arXiv: 1710.05468 [stat.ML].

[LBH15]   Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. "Deep learning". In: *Nature* 521.7553 (May 2015), pp. 436–444. ISSN: 1476-4687. DOI: 10.1038/nature14539. URL: https://doi.org/10.1038/nature14539.

[LL20]   Yulong Lu and Jianfeng Lu. *A Universal Approximation Theorem of Deep Neural Networks for Expressing Probability Distributions*. 2020. arXiv: 2004.08867 [cs.LG].

[Ma+18]   Lei Ma et al. "DeepGauge: Comprehensive and Multi-Granularity Testing Criteria for Gauging the Robustness of Deep Learning Systems". In: *CoRR* abs/1803.07519 (2018). arXiv: 1803.07519. URL: http://arxiv.org/abs/1803.07519.

[Ola]   Christopher Olah. *Neural Networks, Manifolds, and Topology.* https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/. Accessed: 2021-04-29.

[Par+21]   Sejun Park et al. "Minimum Width for Universal Approximation". In: *International Conference on Learning Representations.* 2021. URL: https://openreview.net/forum?id=O-XJwyoIF-k.

[Pei+17]   Kexin Pei et al. "DeepXplore: Automated Whitebox Testing of Deep Learning Systems". In: *CoRR* abs/1705.06640 (2017). arXiv: 1705.06640. URL: http://arxiv.org/abs/1705.06640.

[Qin+20]   Chongli Qin et al. *Training Generative Adversarial Networks by Solving Ordinary Differential Equations.* 2020. arXiv: 2010.15040 [stat.ML].

[RHK18]   Wenjie Ruan, Xiaowei Huang, and Marta Kwiatkowska. *Reachability Analysis of Deep Neural Networks with Provable Guarantees.* 2018. arXiv: 1805.02242 [cs.LG].

[SHK18]   Youcheng Sun, Xiaowei Huang, and Daniel Kroening. "Testing Deep Neural Networks". In: *CoRR* abs/1803.04792 (2018). arXiv: 1803.04792. URL: http://arxiv.org/abs/1803.04792.

[SZ06]   Anton Maximilian Schäfer and Hans Georg Zimmermann. "Recurrent Neural Networks Are Universal Approximators". In: *Proceedings of the 16th International Conference on Artificial Neural Networks - Volume Part I.* ICANN'06. Athens, Greece: Springer-Verlag, 2006, pp. 632–640. ISBN: 3540386254. DOI: 10.1007/11840817_66. URL: https://doi.org/10.1007/11840817_66.

[UM21]   Caterina Urban and Antoine Miné. *A Review of Formal Methods applied to Machine Learning.* 2021. arXiv: 2104.02466 [cs.PL].

[Yos+15]   Jason Yosinski et al. *Understanding Neural Networks Through Deep Visualization.* 2015. arXiv: 1506.06579 [cs.CV].

[Zha+16]   Chiyuan Zhang et al. "Understanding deep learning requires rethinking generalization". In: *CoRR* abs/1611.03530 (2016). arXiv: 1611.03530. URL: http://arxiv.org/abs/1611.03530.

[Zho20]   Ding-Xuan Zhou. "Universality of deep convolutional neural networks". In: *Applied and Computational Harmonic Analysis* 48.2 (2020), pp. 787–794. ISSN: 1063-5203. DOI: https://doi.org/10.1016/j.acha.2019.06.004. URL: https://www.sciencedirect.com/science/article/pii/S1063520318302045.

[Zhu+19]   He Zhu et al. "An Inductive Synthesis Framework for Verifiable Reinforcement Learning". In: *CoRR* abs/1907.07273 (2019). arXiv: 1907.07273. URL: http://arxiv.org/abs/1907.07273.